


Show (well), Attend (better), and Tell (hopefully correctly)

Cheryl Lam, Yanwei C. Liu

Cornell Ann S. Bowers College of Computing and Information Science, Cornell University, Ithaca, NY, 14850



CornellBowers
College of Computing + Information Science

Introduction

Image-captioning requires human-like models that can detect many objects in images and compress salient information into concise, semantically-meaningful captions.

Motivation

- Prior work: *Show and Tell* used a CNN to encode an image into a single vector, then generated a caption with an RNN.
- Limitation: Loses spatial info that could be useful for richer captions.
- *Show, Attend, Tell* introduced **attention** to focus on salient regions of an image during caption generation, improving performance and interpretability.

Goals

- Train and validate caption generators using **soft** and **hard** attention mechanisms on the **Flickr8k** Image-Caption dataset.
- Evaluate the impact of **feature representation** on caption quality.



Human-generated caption:
Professor Weinberger hits a yoda piñata.

Paper Summary

Show, Attend and Tell proposes an attention-based deep neural network model for image caption generation. The model combines a CNN encoder for feature extraction with an LSTM decoder which produces the output caption step-by-step. It introduces an attention mechanism with two variants: 1) soft and 2) hard. The model is evaluated on three datasets of images and captions: Flickr8k, Flickr30k and MS COCO.


Methods



Flickr8k Dataset:

- 8K images
- 5 captions per image

$$e_{ti} = f_{\text{att}}(\mathbf{a}_i, \mathbf{h}_{t-1})$$
$$\alpha_{ti} = \frac{\exp(e_{ti})}{\sum_{k=1}^L \exp(e_{tk})}$$



Soft Attention

$$\mathbb{E}_{p(s_t|a)}[\hat{\mathbf{z}}_t] = \sum_{i=1}^L \alpha_{t,i} \mathbf{a}_i$$

Doubly Stochastic Attention

$$L_d = -\log(P(\mathbf{y}|\mathbf{x})) + \lambda \sum_i^L (1 - \sum_t^C \alpha_{ti})^2$$

Hard Attention

$$\tilde{s}_t \sim \text{Multinoulli}_L(\{\alpha_i\})$$
$$p(s_{t,i} = 1 | s_{j<t}, \mathbf{a}) = \alpha_{t,i}$$
$$\hat{\mathbf{z}}_t = \sum_i s_{t,i} \mathbf{a}_i.$$

Trained using REINFORCE

What does the model “see”?



A dog is playing in the grass. ✓



A man in a blue shirt is standing on a rock overlooking a snowy mountain. ✓




A boy in a blue shirt is riding a bike on a tire swing. ✗

Results


Model	BLEU				METEOR
	BLEU-1	BLEU-2	BLEU-3	BLEU-4	
Soft-Attention VGG (Xu, K., et al., 2015)	67.00	44.80	29.90	19.50	18.93
Soft-Attention VGG	65.67	45.67	32.14	22.61	20.72
Soft-Attention ResNet	68.01	48.30	34.39	24.20	21.66
Hard-Attention VGG (Xu, K., et al., 2015)	67.00	45.70	31.40	19.50	20.30
Hard-Attention VGG	64.72	43.57	29.87	20.46	19.28
Hard-Attention ResNet	68.31	48.75	34.67	24.40	22.04

Table 1. BLEU-1,2,3,4/METEOR metrics on Flickr8k Dataset.

Soft Attention



Hard Attention



Conclusion

- Attention improves quality/interpretability and we outperformed paper results.
- **Stronger feature extractors (ResNet50) produce better captions than VGG19.**
- Hard attention is harder to train than soft attention due to its stochastic nature.

References

Xu, K., Ba, J., Kiros, R., Cho, K., Courville, A., Salakhudinov, R., ... & Bengio, Y. (2015, June). Show, attend and tell: Neural image caption generation with visual attention. In *International conference on machine learning* (pp. 2048-2057). PMLR.

Hodosh, M., Young, P., & Hockenmaier, J. (2013). Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47, 853–899.

We would like to thank Professor Weinberger, Professor Ma, and the CS 4782 Course Staff.